

The Brutal Q&A

Every hard question about Inayat — in points you can actually remember

No passages to re-read. Short points, in order. Say the bullets in your own words — don't recite them.

0. The name — say this before anything else

Q: What does “Inayat” mean, and why did you choose it?

- Inayat means “care” in Urdu.
- We named it that because the whole prototype exists to give every patient the caring attention a rushed two-minute doctor visit can't.

Q: Which dataset did you train it on?

- None — we didn't train a model at all.
- The danger-check and follow-up questions are fixed rules, not learned from data.
- The “AI” part is a real, already-trained model (Claude) doing live reasoning — not something we trained ourselves.
- Our only “data” is a small, 14-sentence reference list we wrote ourselves, and openly say isn't officially certified yet.

Q: What languages does it actually support?

- English, Hindi, and Telugu — both on-screen text and voice.
- Picked Telugu specifically since this problem statement's example district (Kurnool) is Telugu-speaking.

Q: Does it work with no internet at all?

- The page itself still needs one internet load, like any website.
- But once loaded: if the AI can't be reached, it doesn't crash or freeze.
- It safely falls back to “route this to a human,” instead of guessing.

Q: Can a patient ask you to delete their data?

- Not yet, honestly — there's no delete button today.
- A doctor can correct or amend a case, but records aren't destroyed.
- A real, known gap, not something we're hiding.

Q: Is the code actually public? Can we check it ourselves?

- Yes — real GitHub repo, real commit history, nothing hidden.
 - github.com/azlanabyssal-cloud/carepilot
-

1. The problem — from zero, no background needed

Picture this first, before any numbers: someone is sick, waits hours in a government hospital line, and gets a doctor who has only minutes before the next patient. That's the whole problem, in one picture. The numbers below just prove it's real and everywhere.

Q: What is the actual problem?

- A government-hospital doctor sees a patient for about **2 minutes**.
- Not enough time to ask, listen, AND read old prescriptions.
- The “listening” part is what gets rushed first.
- But real studies say: careful listening is the #1 reason doctors get the diagnosis right.

Q: Why does this matter for real people, not just in theory?

- 2 out of every 3 Indians live in villages.
- Only about 1 in 4 doctors work in villages.
- Least access to a doctor = also the least time once you get one.

Q: How did you solve it, in one sentence?

- We talk to the patient **before** the doctor does.
 - Type, speak in their own language, or photo an old prescription.
 - We turn that into one clean, organized note.
 - By the time the doctor walks in, the listening is already done.
-

2. The pipeline — 7 steps, and why each one exists

- **1. Consent.** Patient must agree first — system refuses to save without it.
 - **2. Patient talks.** Type, speak, or photo — not everyone types well or speaks English.
 - **3. Smart follow-up questions.** Cough → phlegm/colour questions. Rash → spreading/new-soap questions. Pain → the standard 8-question method doctors already use. One list for everything would be lazy.
 - **4. Danger check FIRST.** Checks exact words against a dangerous-symptom list. Zero AI needed. Zero internet needed. A dangerous case should never depend on AI having a good day.
 - **5. AI helps (if not obviously dangerous).** Suggests urgency. No AI available? A safe backup kicks in instead of the app failing.
 - **6. Second, independent double-check.** Compares words against real medical guidelines. Can only push urgency UP, never down. A wrong “less serious” correction is far more dangerous than extra caution.
 - **7. Doctor decides.** Everything saved. Doctor reads it, can change anything. AI never diagnoses. AI never prescribes.
-

3. The engineering questions

Q: How much space (GB), and where does it run?

- Real app code: about **4.5 MB** — smaller than one phone photo.
- Libraries it actually needs to run: about **290 MB**, measured directly.
- One heavy library (PyTorch, for a feature not wired into the live app yet) left OUT on purpose — saves ~5 GB of dead weight.
- Hosted on **Render**, a cloud host, free tier.
- Free tier sleeps after 15 min idle, takes 30–60 sec to wake up — we say this openly.

Q: What parameters or choices did you make, and why?

- Guideline matching: simple word-matching, not a giant AI model.
- Why: our reference list is only a few dozen short passages — a big model gains nothing here.
- We trust only the **single best** guideline match, not the top 3.
- Why: a weak 3rd-place match once wrongly escalated “mild knee pain” to EMERGENCY.
- We caught it ourselves, tested it, fixed it.

Q: What tech did you actually build this with?

- Python for the server.
 - Plain HTML/JavaScript for the screen — no heavy framework, loads fast on cheap phones.
 - A small database to save cases.
 - Claude (AI model) for language understanding.
 - Bhashini — real Indian-language voice tool.
 - ABDM — India's real health-ID system.
-

4. Why we're different — and the ONE thing nobody else has

Q: How is this better than another team's health chatbot?

- Most chatbots: one AI model + a disclaimer.
- Ours: a safety net that doesn't need the AI at all.
- Keeps working with zero internet.
- Saves the case for the doctor to reopen later.
- Connects to real government health systems.

Q: The ONE thing you have that basically nobody else has?

- A safety check that can only get MORE careful, never less.
- Works even if the AI is completely wrong, or fully offline.
- Most student projects trust the AI's answer directly.
- We built one that never fully trusts it.

Q: Isn't this just ChatGPT with extra steps?

- ChatGPT: no double-check, nothing saved, stops with no internet.
 - We built 3 separate real things ChatGPT doesn't have.
 - Not one clever prompt — real engineering.
-

5. Grabbing attention — the pitch, and the app itself

Q: How will you grab attention when you present this?

- Open with the real number, not a claim.
- “2 minutes per patient” — one of the lowest in the world.
- From a real study of 67 countries.
- Everyone's sat in that waiting room — it lands instantly.

Q: Once someone is looking at the app, what makes IT grab them?

- Homepage doesn't wait for a click.
 - A real example types itself out live, on its own.
 - Gets caught as an emergency — before the visitor touches anything.
 - Proof, not a promise.
 - Page reveals content smoothly while scrolling — feels alive, not static.
-

6. Limitations — proof we KNOW them, not guessed

A real limitation vs. a guess: we tested it and watched it fail, with our own eyes, before writing it down.

- **Old handwritten prescriptions read worse than clean typed ones.** Proof: same code, blurry image turned “13.0-17.0” into wrong “130-170.” We watched it happen.
 - **Our medical reference list isn't officially certified yet.** We wrote it ourselves as a real starting point — we say so.
 - **Voice and health-ID features never touched real government servers.** No test credentials exist for us — only public documentation.
 - **Offline voice recognition isn't very accurate yet.** We tested and measured it ourselves — every result gets flagged “needs a human check,” never silently trusted.
-

7. The extra brutal round

Q: What's the hardest real bug you personally found and fixed?

- Safety-check used to trust the “most serious” of the top 3 guideline matches.
- A weak, barely-related 3rd match once beat a correct, strong 1st match.
- Wrongly turned ordinary “mild knee pain” into EMERGENCY.
- We found it ourselves, proved it with a test.
- Fixed: only trust the single best match now.

Q: What if the AI gives a wrong answer during the demo?

- Danger-word check runs FIRST — zero AI needed for that.
- Second guideline check can only make things MORE careful.
- One wrong AI answer alone still can't create an unsafe result.
- Would need TWO independent systems wrong, in the same direction, at once.

Q: Couldn't a big company clone this in a week?

- The AI-chatbot part? Maybe, yes.
- The safety-net design, real integrations, real-world testing? No.
- That's not a weekend job — most teams skip it entirely.

Q: Is patient data actually safe?

- Doctor's screen sits behind a real login.
- Wrong passcode = the server itself refuses the data.
- Not just a hidden button someone could work around.

Q: Have you tested this on real patients in a real hospital?

- No — and we say that honestly.
- 367 automated tests, all currently passing.
- Real example cases we wrote ourselves.
- Real-patient testing is a genuine next step, not a claim we're making.

Q: If teammates give slightly different answers on stage, is that bad?

- No — only guessing looks bad.
 - Slightly different wording is fine.
 - Don't know something? Say so, offer to follow up.
-

8. The answers you say without thinking

The whole thing, in one line: “Inayat listens to a patient before the doctor does, and turns what they say into a safe, organized note — so the doctor's few minutes go to deciding, not repeating questions.”

Three reasons people should respect this project:

- We say what's NOT finished, out loud, before you catch us.
- The safety design doesn't trust one single thing to be right — it catches itself being wrong.
- Every real government feature we claim, we actually built and tested.

You've now read the real problem, the real pipeline, the real numbers, the real bugs we found ourselves, and the real limits. That's not luck. That's understanding. Go say it.