

Inayat — SIH26047 Pitch Content Guide

What to put in the PPT, what to leave out, and exactly why — every line below checked against the live code and the real problem-statement text on 13 Sep 2026, not assumed.

This is not a features list. It's a filter: for every claim you might want to put on a slide, this tells you whether the code actually backs it up, and if not, what the honest version of that claim is. A judge who catches one overclaim discounts everything else on the slide with it — this document exists so that never happens.

1. PS Alignment Scorecard

SIH26047 ("Patient Case-Taking Software", Ministry of Ayush / All India Institute of Ayurveda) names four modules. Here is what's actually built against each, not what was planned.

PS Module	What's Real	PPT Guidance
A — Conversational Multimodal History Engine	Real SOCRATES-framework follow-up (8 standard questions), plus real category-specific templates for respiratory/GI/dermatological/general complaints — deterministic, zero-LLM, clinically standard. Now a real one-question-at-a-time conversation with skip controls, not a static list. Red-flag detection is a real deterministic keyword scanner, verified live.	Claim it directly. Say "dialogue manager constrained by a clinical ontology," not "AI chatbot" — that's literally what the PS itself asks for and what this is.
B — Medical Document Digitization	Real OCR (Tesseract) + chronological multi-document ordering + abnormal lab-value flagging, all live-tested. Medication extraction is a regex heuristic (drug-name-shaped word next to a dosage number) — can mismatch on real documents. Drug-interaction flagging (a distinct PS bullet from lab-value flagging) is not built. Handwriting OCR only tested against clean synthetic images, never real messy handwriting.	Claim the timeline + lab-value flagging. Do not claim "structured medication extraction" or "handwriting OCR" without a live test on a real messy sample first.
C — Structured History Summary Generator	Real Chief Complaint → HPI → Past/Drug/Family/Personal/ROS structure exists. Physician edit/confirm gate is real and enforced. But: without a live Anthropic API key, every summary is a template sentence built from raw fields, not a real synthesized narrative — fixed to read naturally this session, but still not clinical reasoning.	Only claim "AI-drafted summary" if a real key is configured for the demo. Otherwise say "structured intake with a safety-net fallback," which is the honest, still-impressive truth.
D — Consent, Privacy & ABDM Integration	Real, enforced consent gate (client + server side). Real ABDM sandbox adapter — genuine RSA-OAEP encryption matching ABDM's published spec, 18 passing tests. But: this is OTP/enrollment against the ABDM sandbox, not a live push to a hospital HIS/EMR via FHIR.	Claim "real ABDM sandbox integration, cryptographically correct." Do not claim "integrated with the hospital's HIS" — that link doesn't exist.

AYUSH mode (Dashavidha Pariksha) doesn't map to one module — see Section 2.

2. Lead With These — Genuinely Strong, Verified

Asymmetric escalate-only safety design

The Guideline-Verification agent can only raise a triage level, never lower one. This is a real, correct clinical-safety pattern — say this explicitly in the pitch, it's the single most defensible design decision in the system. It answers the "what if your AI is wrong" question before a judge asks it.

Zero-API resilience, proven live

Offline speech recognition (PocketSphinx, English), offline text-to-speech (espeak-ng, English/Hindi/Telugu), and deterministic reasoning/history-drafting fallbacks all verified working with zero credentials and zero network calls. Most student prototypes have a single point of failure at their API key — this one degrades honestly instead of breaking.

An evaluation number that's actually honest

The safety-metrics card shows a raw fraction (e.g. "4/4"), not a percentage computed from a handful of cases — fixed this session specifically because a bare "100%" next to a small sample reads as either not understanding statistics or hiding the fine print. Say the number out loud as a fraction in the pitch, and say plainly what it does and doesn't prove.

AYUSH parameter split, with real citations

The Dashavidha Pariksha implementation deliberately separates patient-self-report parameters (Prakriti, Vikriti, Satmya, Sattva, Ahara Shakti, Vyayama Shakti, Vaya) from physician-only-assessment parameters (Sara, Samhanana, Pramana — traditionally assessed by physical examination, not self-report), each gloss cross-checked against published sources including CCRAS (Ministry of Ayush's own research body). This is the opposite of "decorative Sanskrit vocabulary" — it shows real respect for what a kiosk can and can't legitimately ask a patient. Worth its own slide, especially for an AIIA/Ministry-of-Ayush audience specifically.

Real cryptography, not a mocked integration

The ABDM adapter implements actual RSA-OAEP encryption matching ABDM's published certificate spec (SHA-1/MGF1 padding), with 18 passing tests including a real encrypt/decrypt round-trip. If a judge asks "is this integration real or just a UI mockup," the honest answer is "the cryptography is real and tested; the OTP flow talks to ABDM's real sandbox."

Multilingual end to end, not just the UI chrome

English/Hindi/Telugu isn't just button labels — voice input, voice output, and the case summary itself all work in all three languages, verified live this session, including the offline fallback paths. This is a real, judge-visible signal of contextual awareness for an India-facing OPD tool — don't undersell it to one line on a slide.

3. Do Not Claim These — Or Only With The Honest Wording Given

"AI-powered clinical reasoning" — only if a real API key is live for the demo

Without ANTHROPIC_API_KEY configured, every summary comes from a deterministic template. It was rewritten this session to read as one real sentence instead of three bolted-together log lines, but it is still not clinical reasoning. If the key isn't set before demo day, say "structured intake with a zero-API safety net" — true, and still a real differentiator (most teams have no fallback at all).

"Clinical reasoning engine" for the guideline check — it's text similarity

The Guideline-Verification agent uses TF-IDF cosine similarity against a guideline corpus — real, working, and asymmetrically safe (see Section 2), but it is lexical text matching, not structured symptom-cluster reasoning. It will do well on textbook-phrased complaints and can miss atypical presentations (the same population where triage tools historically fail worst). If a judge asks about atypical presentations, the honest, still-credible answer is: "the asymmetric design means a weak match never lowers your priority, but the match itself is lexical — that's a named next step, not a hidden gap."

"Handwriting OCR" — unverified

Only tested against clean, printed synthetic images. Real handwritten prescriptions have not been tried. Do not put a handwritten prescription in the live demo unless you've tested that exact image beforehand and it actually works.

"Structured medication extraction" — it's a regex heuristic

The extractor looks for a capitalized word next to a dosage-shaped number. It can grab a dosage-form word ("Tablet 500mg") instead of the actual drug name when the layout doesn't match its assumption. Frame this as "candidate extraction for physician confirmation," not "AI drug recognition."

"Integrated with ABDM / the hospital HIS" — it's sandbox-only

The real, tested part is OTP request/confirm against ABDM's sandbox with correct encryption. There is no live push of a structured summary into a hospital's HIS or FHIR store. Say "ABDM sandbox-integrated," not "HIS-integrated."

Any percentage without its sample size next to it

The instinct to fix this once (Section 2) should generalize: any number on any slide — accuracy, recall, time saved — needs its real sample size or basis stated in the same breath, not in footnote text a judge has to search for.

4. Slide-by-Slide

Slide	Put This	Not This
Problem	The 2-5 minute OPD consult stat (BMJ Open, 2017 — real, cited); the fragmented-records problem; AYUSH's extra history-taking burden (Dashavidha vs allopathic).	Generic "healthcare is broken" framing that doesn't cite this PS's own named numbers.
Solution Overview	The 5-step patient journey exactly as the PS names it (Identify → Converse → Scan → Summarize & Route → Consult) mapped one-to-one to what's actually built.	A generic "AI does everything" diagram that doesn't correspond to real endpoints.
Architecture	The zero-API fallback chain (real backend preferred, deterministic safety net always available) — this is your actual technical differentiator versus a wrapper-around-an-API-call competitor.	A slide that only shows the happy-path API call and hides what happens when it's unavailable.
Live Demo	Run it live at the deployed URL. Show the SOCRATES conversation branching by complaint type (cough → respiratory questions, not generic pain questions) — this is real, verified, and visibly "smart" without needing an API key.	A pre-recorded video as the only "proof" — judges discount video demos versus a live one that can be redirected.
Impact / Validation	The safety-metrics fraction, stated with its real sample size out loud, plus the asymmetric-escalation design as your safety argument.	Any bare percentage. Any accuracy claim without "measured against an N-case, author-labeled set — not clinical validation."
Roadmap	Structured feature extraction to reduce reliance on text-similarity for guideline matching; drug-interaction flagging; real handwriting-OCR validation; production HIS/FHIR push.	Presenting any of these as already done.

5. Priority Order Before Demo Day

- **Decide on the Anthropic API key.** This single decision determines whether Slide "Solution Overview" can say "AI-drafted" or must say "structured intake with a safety net." Decide first — everything else about wording follows from it.
- **Test the live demo path end to end, today, on the actual deployed URL** — not localhost. Confirm the emergency red-flag case, a normal case through the full SOCRATES conversation, and the physician console login all work on a phone, on real mobile data, not just office wifi.
- **Rehearse the three honest answers** from Section 3 (atypical presentations, medication extraction, ABDM scope) so they come out as confident, prepared answers, not backpedaling.
- **Do not add AYUSH/ABDM toggles to the emergency-case demo path** — already fixed in code to hide automatically, but make sure your live click-through doesn't manually surface them there.
- **If time allows, get one real (not synthetic) test of a messy handwritten prescription** — either confirm it works enough to show, or plan the demo around a clean printed one instead.

Every claim above was checked against the actual repository and a live run of the deployed application on 13 Sep 2026 — grep results, direct API calls, and Playwright-driven browser tests, not memory or assumption. If the code changes after this date, re-verify before trusting this document again.